

# Enhancing and Evaluating Probabilistic Circuits for High-Resolution Lossless Image Compression

Daniel Severo\*, Jingtong Su<sup>†</sup>, Anji Liu<sup>§</sup>, Jeff Johnson\*, Brian Karrer\*  
Guy Van den Broeck<sup>§</sup>, Matthew J. Muckley\*, Karen Ullrich\*

\*Meta AI

<sup>†</sup>NYU

<sup>§</sup>UCLA

`{dsevero,jhj,briankarrer,mmuckley,karenu}@meta.com`

`js12196@nyu.edu`

`{liuanji,guyvdb}@cs.ucla.edu`

# Lossless Neural Compression

# Lossless Neural Compression

## Lossless compression

- density/PMF estimation
- entropy coding (ANS, AC)

# Lossless Neural Compression

## Lossless compression

- density/PMF estimation
- entropy coding (ANS, AC)

Entropy coding  $X_1, \dots, X_n$  requires computing  $P(X_i < x | X_1 = x_1, \dots, X_{i-1} = x_{i-1})$

# Lossless Neural Compression

## Lossless compression

- density/PMF estimation
- entropy coding (ANS, AC)

Entropy coding  $X_1, \dots, X_n$  requires computing  $P(X_i < x | X_1 = x_1, \dots, X_{i-1} = x_{i-1})$

Lossless Neural Compression = Model P with a neural network

# Lossless Neural Compression

At least 2 ways to define a model for  $P_{\theta}(X_1, \dots, X_n)$

# Lossless Neural Compression

At least 2 ways to define a model for  $P_{\theta}(X_1, \dots, X_n)$

Bottom-up:  $P_{\theta_i}(X_i | X_1, \dots, X_{i-1})$

# Lossless Neural Compression

At least 2 ways to define a model for  $P_{\theta}(X_1, \dots, X_n)$

Bottom-up:  $P_{\theta_i}(X_i | X_1, \dots, X_{i-1}) \Rightarrow \prod_{i=1}^n P_{\theta_i}(X_i | X_1, \dots, X_{i-1}) = P_{\theta}(X_1, \dots, X_n)$



# Lossless Neural Compression

At least 2 ways to define a model for  $P_\theta(X_1, \dots, X_n)$

Bottom-up:  $P_{\theta_i}(X_i | X_1, \dots, X_{i-1}) \Rightarrow \prod_{i=1}^n P_{\theta_i}(X_i | X_1, \dots, X_{i-1}) = P_\theta(X_1, \dots, X_n)$

Top-down:  $P_\theta(X_1, \dots, X_n)$

# Lossless Neural Compression

At least 2 ways to define a model for  $P_\theta(X_1, \dots, X_n)$

Bottom-up:  $P_{\theta_i}(X_i | X_1, \dots, X_{i-1}) \Rightarrow \prod_{i=1}^n P_{\theta_i}(X_i | X_1, \dots, X_{i-1}) = P_\theta(X_1, \dots, X_n)$

Top-down:  $P_\theta(X_1, \dots, X_n) \Rightarrow \text{marginalization} \Rightarrow P_{\theta_i}(X_i | X_1, \dots, X_{i-1})$

# Lossless Neural Compression

At least 2 ways to define a model for  $P_{\theta}(X_1, \dots, X_n)$

Bottom-up:  $P_{\theta_i}(X_i | X_1, \dots, X_{i-1}) \Rightarrow \prod_{i=1}^n P_{\theta_i}(X_i | X_1, \dots, X_{i-1}) = P_{\theta}(X_1, \dots, X_n)$

Top-down:  $P_{\theta}(X_1, \dots, X_n) \Rightarrow \text{marginalization} \Rightarrow P_{\theta_i}(X_i | X_1, \dots, X_{i-1})$

**Probabilistic Circuits are top-down models with efficient marginalization properties**

# Chow-Liu Trees (Chow & Liu, 1968)

# Chow-Liu Trees (Chow & Liu, 1968)

Given a joint  $P(X_1, \dots, X_n)$ , what's the best model

we can construct from products of conditionals

$$Q(X_1, \dots, X_n) = \prod_i P(X_i | X_{j(i)}) \quad ?$$

# Chow-Liu Trees (Chow & Liu, 1968)

Given a joint  $P(X_1, \dots, X_n)$ , what's the best model

we can construct from products of conditionals

$$Q(X_1, \dots, X_n) = \prod_i P(X_i | X_{j(i)}) \quad ?$$

Goodness is measured by  $\text{KL}[P||Q]$

# Chow-Liu Trees (Chow & Liu, 1968)

Given a joint  $P(X_1, \dots, X_n)$ , what's the best model

we can construct from products of conditionals

$$Q(X_1, \dots, X_n) = \prod_i P(X_i | X_{j(i)}) \quad ?$$

Goodness is measured by  $\text{KL}[P||Q]$

The only optimization variable is the assignment

function  $j: \{1, \dots, n\} \mapsto \{1, \dots, n\}$


# Chow-Liu Trees (Chow & Liu, 1968)

Given a joint  $P(X_1, \dots, X_n)$ , what's the best model we can construct from products of conditionals

$$Q(X_1, \dots, X_n) = \prod_i P(X_i | X_{j(i)}) \quad ?$$

Goodness is measured by  $\text{KL}[P||Q]$

The only optimization variable is the assignment function  $j: \{1, \dots, n\} \mapsto \{1, \dots, n\}$



This is equivalent to asking which tree-shaped probabilistic graphical model minimizes the KL divergence?



# Chow-Liu Trees (Chow & Liu, 1968)

If  $Q$  is a tree, then it can be shown that

$$\text{KL} [P||Q] = - \sum_i I(X_i; X_{j(i)}) + \text{const}$$

# Chow-Liu Trees (Chow & Liu, 1968)

If  $Q$  is a tree, then it can be shown that

$$\text{KL}[P||Q] = - \sum_i I(X_i; X_{j(i)}) + \text{const}$$

Chow-Liu Tree Algorithm:

- 1) Build a complete weighted graph with MI as edge weights

# Chow-Liu Trees (Chow & Liu, 1968)

If  $Q$  is a tree, then it can be shown that

$$\text{KL}[P||Q] = - \sum_i I(X_i; X_{j(i)}) + \text{const}$$

Chow-Liu Tree Algorithm:

- 1) Build a complete weighted graph with MI as edge weights
- 2) Find a **maximum spanning tree (MST)**




# Chow-Liu Trees (Chow & Liu, 1968)

If  $Q$  is a tree, then it can be shown that

$$\text{KL}[P||Q] = - \sum_i I(X_i; X_{j(i)}) + \text{const}$$

Chow-Liu Tree Algorithm:

- 1) Build a complete weighted graph with MI as edge weights
- 2) Find a **maximum spanning tree (MST)**



# Chow-Liu Trees (this part isn't in the original paper)

If  $P$  is not known ....

# Chow-Liu Trees (this part isn't in the original paper)

If  $P$  is not known ....., then it can be shown that

$$\text{KL} [P||Q] = \underbrace{-\sum_i I(P_{X_i, X_{J(i)}})}_{\text{Original CLT objective (function of } J \text{ only)}} + \Delta H(P_X)$$

# Chow-Liu Trees (this part isn't in the original paper)

If  $P$  is not known ..., then it can be shown that

$$\text{KL} [P||Q] = \underbrace{-\sum_i I(P_{X_i, X_{J(i)}}) + \Delta H(P_X)}_{\text{Original CLT objective (function of } J \text{ only)}} + \underbrace{\sum_i \mathbb{E}_{X_{J(i)}} \text{KL} \left[ P_{X_i | X_{J(i)}}(\cdot | X_{J(i)}) || Q_{X_i | X_{J(i)}}^{\theta_i, J}(\cdot | X_{J(i)}) \right]}_{\text{Estimation error (function of both } J, \theta)},$$

where  $\Delta H(P_X) = \sum_i H(P_{X_i}) - H(P_X) \geq 0$  with equality when  $X_i$  are pair-wise independent.

# Chow-Liu Trees (this part isn't in the original paper)

If  $P$  is not known ..., then it can be shown that

$$\text{KL} [P||Q] = \underbrace{-\sum_i I(P_{X_i, X_{J(i)}}) + \Delta H(P_X)}_{\text{Original CLT objective (function of } J \text{ only)}} + \underbrace{\sum_i \mathbb{E}_{X_{J(i)}} \text{KL} \left[ P_{X_i | X_{J(i)}}(\cdot | X_{J(i)}) || Q_{X_i | X_{J(i)}}^{\theta_i, J}(\cdot | X_{J(i)}) \right]}_{\text{Estimation error (function of both } J, \theta)},$$

where  $\Delta H(P_X) = \sum_i H(P_{X_i}) - H(P_X) \geq 0$  with equality when  $X_i$  are pair-wise independent.

Joint optimization over  $(\theta, J)$  is difficult



# Chow-Liu Trees (this part isn't in the original paper)

If  $P$  is not known ....., then it can be shown that

$$\text{KL} [P||Q] = \underbrace{-\sum_i I(P_{X_i, X_{J(i)}}) + \Delta H(P_X)}_{\text{Original CLT objective (function of } J \text{ only)}} + \underbrace{\sum_i \mathbb{E}_{X_{J(i)}} \text{KL} \left[ P_{X_i | X_{J(i)}}(\cdot | X_{J(i)}) || Q_{X_i | X_{J(i)}}^{\theta_i, J}(\cdot | X_{J(i)}) \right]}_{\text{Estimation error (function of both } J, \theta)},$$

where  $\Delta H(P_X) = \sum_i H(P_{X_i}) - H(P_X) \geq 0$  with equality when  $X_i$  are pair-wise independent.

Joint optimization over  $(\theta, J)$  is difficult

Pragmatic solution: two-step process

# Chow-Liu Trees (this part isn't in the original paper)

If  $P$  is not known ....., then it can be shown that

$$\text{KL} [P||Q] = \underbrace{-\sum_i I(P_{X_i, X_{J(i)}}) + \Delta H(P_X)}_{\text{Original CLT objective (function of } J \text{ only)}} + \underbrace{\sum_i \mathbb{E}_{X_{J(i)}} \text{KL} \left[ P_{X_i | X_{J(i)}}(\cdot | X_{J(i)}) || Q_{X_i | X_{J(i)}}^{\theta_i, J}(\cdot | X_{J(i)}) \right]}_{\text{Estimation error (function of both } J, \theta)},$$

where  $\Delta H(P_X) = \sum_i H(P_{X_i}) - H(P_X) \geq 0$  with equality when  $X_i$  are pair-wise independent.

Joint optimization over  $(\theta, J)$  is difficult

Pragmatic solution: two-step process

- Estimate the MI of every pair  $I(X_i; X_j)$ , then find the MST (equivalently,  $J^*$ )

# Chow-Liu Trees (this part isn't in the original paper)

If  $P$  is not known ....., then it can be shown that

$$\text{KL} [P||Q] = \underbrace{-\sum_i I(P_{X_i, X_{J(i)}}) + \Delta H(P_X)}_{\text{Original CLT objective (function of } J \text{ only)}} + \underbrace{\sum_i \mathbb{E}_{X_{J(i)}} \text{KL} \left[ P_{X_i | X_{J(i)}}(\cdot | X_{J(i)}) || Q_{X_i | X_{J(i)}}^{\theta_i, J}(\cdot | X_{J(i)}) \right]}_{\text{Estimation error (function of both } J, \theta)},$$

where  $\Delta H(P_X) = \sum_i H(P_{X_i}) - H(P_X) \geq 0$  with equality when  $X_i$  are pair-wise independent.




Joint optimization over  $(\theta, J)$  is difficult

Pragmatic solution: two-step process





- Estimate the MI of every pair  $I(X_i; X_j)$ , then find the MST (equivalently,  $J^*$ )
- Minimize estimation error by optimizing  $Q_{X_i | X_{J(i)}}^{\theta_i, J}$

# Hidden CLTs (Liu & Van den Broeck, 2021)

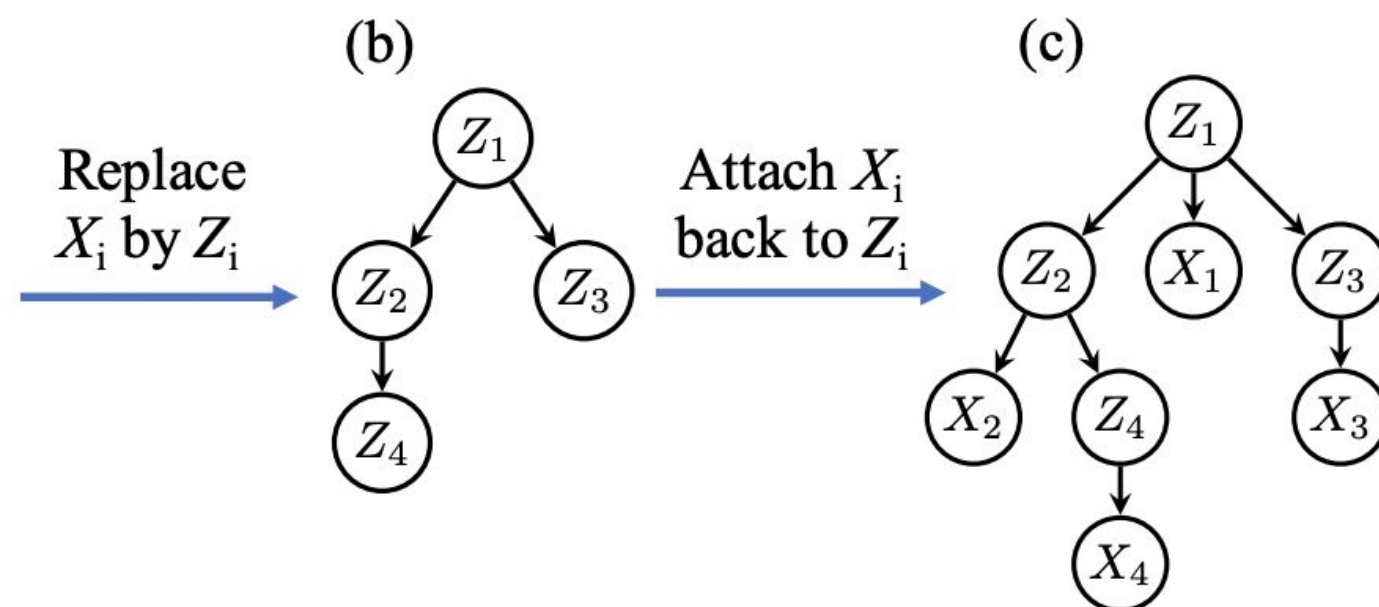



# Hidden CLTs (Liu & Van den Broeck, 2021)



# Hidden CLTs (Liu & Van den Broeck, 2021)

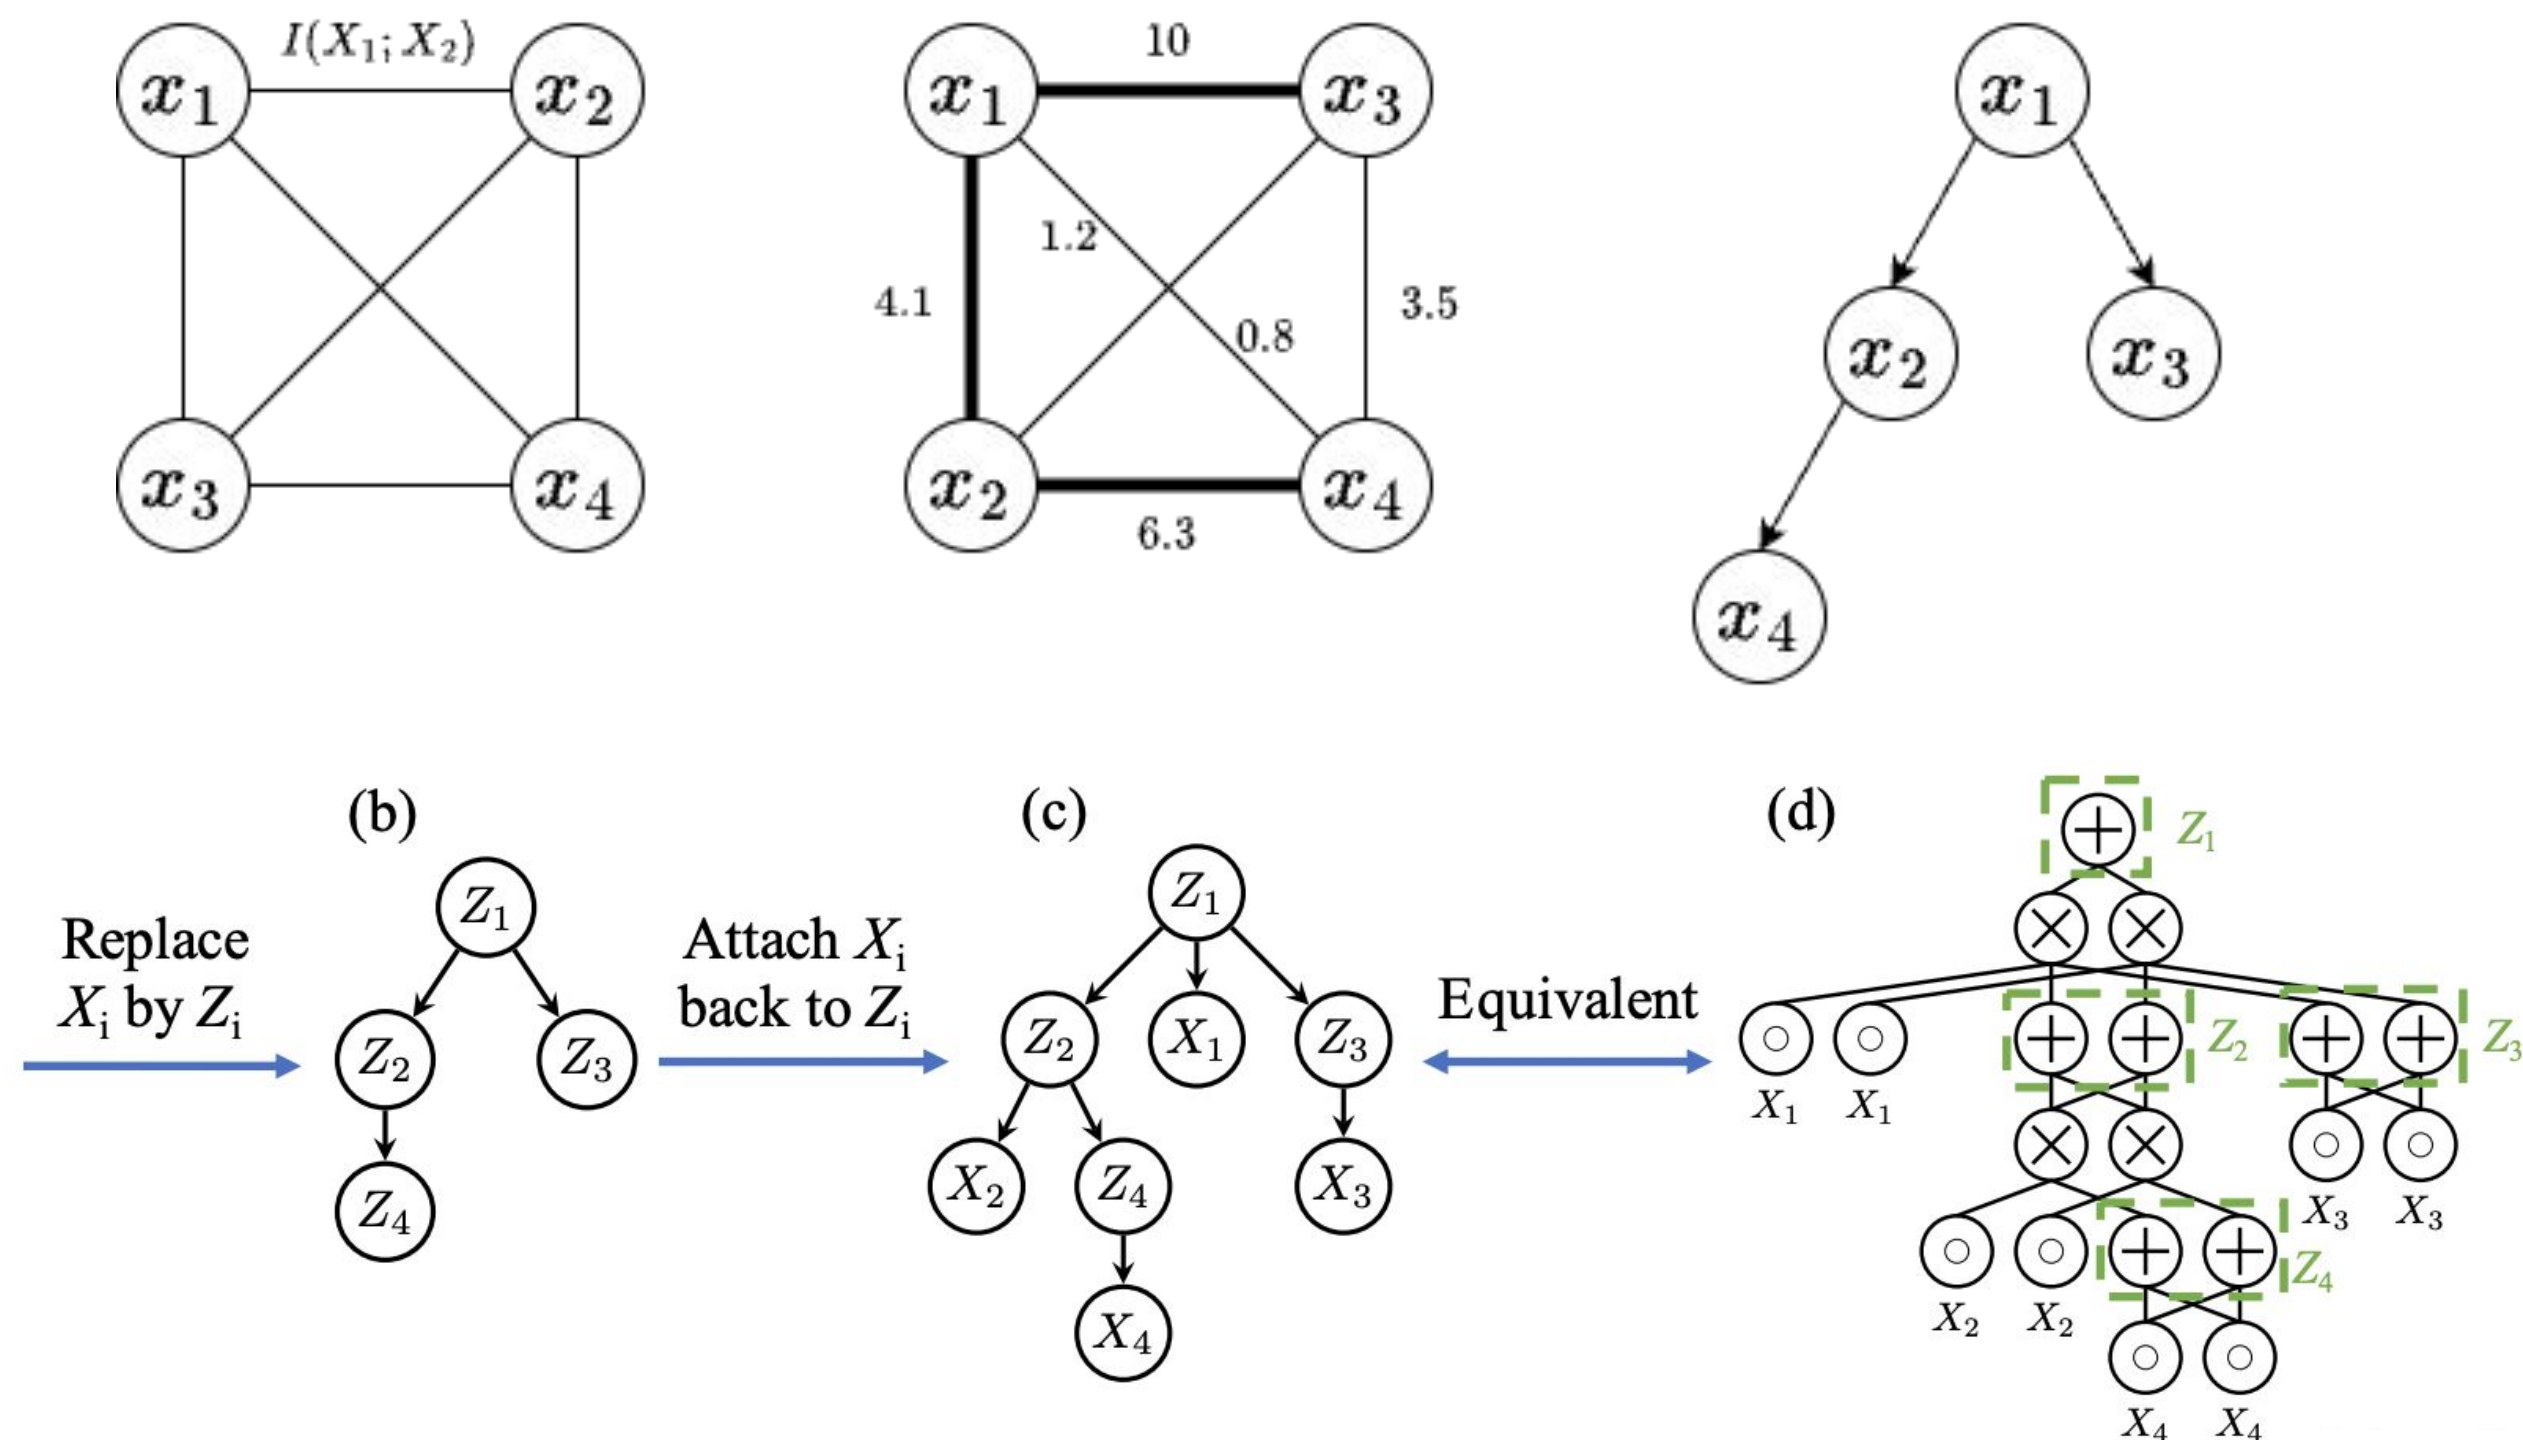


# Hidden CLTs (Liu & Van den Broeck, 2021)





# Hidden CLTs (Liu & Van den Broeck, 2021)





# Our contributions

# Sparse Mutual Information Estimation

Estimating MI scales quadratically with  $n$

# Sparse Mutual Information Estimation

Estimating MI scales quadratically with  $n$

For images, MI will correlate with pixel distance

# Sparse Mutual Information Estimation

Estimating MI scales quadratically with  $n$

For images, MI will correlate with pixel distance


**Solution:** only compute MI between neighbouring pixels, set others to 0

# Sparse Mutual Information Estimation

Estimating MI scales quadratically with  $n$


For images, MI will correlate with pixel distance

**Solution:** only compute MI between neighbouring pixels, set others to 0




# Repurpose Invertible Transforms from WebP

ConvertColorSpace (RGB or YCoCg)



PredictTransformThenConvertColorspacePipeline (YCoCg)




Only improved performance of PCs, not  
other neural compressors we tried!



# Scaling up to AR/VR avatar datasets

Dataset to train 3D renderings of people created by Codec Avatar team at Meta

- [publicly released on github](#).
- 50TB uncompressed images
- resolution of  $2048 \times 1334$  pixels





# Experiments

Table 1: PC results compared to standard codecs and neural compressors. PCs provide a gain in compression performance relative to standard codecs, while not requiring as much compute as neural compressors. The advantage of both PCs and neural codecs diminishes as the average image size of the dataset increases (left-to-right). All units are bits-per-dimension.

		CIFAR	IM32	IM64	CLIC	MF
Standard Codecs	PNG(RGB)	5.87	6.05	5.34	3.49	2.87
	PNG(YCoCg)	5.23	5.54	4.88	3.13	3.01
	WebP(YCoCg)	4.87	5.20	4.51	2.68	2.66
	WebP(RGB)	<b>4.61</b>	<b>4.98</b>	<b>4.30</b>	<b>2.59</b>	<b>2.61</b>
PCs	HCLT	6.04	6.16	5.92	4.10	3.12
	HCLT++	<b>4.13</b>	<b>4.72</b>	<b>4.29</b>	<b>2.48</b>	<b>2.75</b>
Neural Codecs	HiLLoC	3.56	4.20	3.90	2.63	-
	IDF	3.32	3.95	3.66	2.43	2.57
	IDF++	<b>3.26</b>	<b>3.94</b>	<b>3.62</b>	<b>2.44</b>	<b>2.54</b>
Adv. of best Prob. Circuit over WebP		10.8%	5.8%	1.0%	4.0%	-5.4%
Adv. of best Neural Codec over WebP		29.2%	20.8%	15.9%	6.1%	2.7%



# Obrigado! / Thanks!

Slides will be available at [dsevero.com](https://dsevero.com)

# References

[1] (2020) Probabilistic Circuits: A Unifying Framework for Tractable Probabilistic Models

<https://yoojungchoi.github.io/files/ProbCirc20.pdf>

[2] (2022) Lossless Compression with Probabilistic Circuits

<https://arxiv.org/pdf/2111.11632.pdf>

[3] (2020) Group Fairness by Probabilistic Modeling with Latent Fair Decisions

<https://arxiv.org/pdf/2009.09031.pdf>

This paper outlines the full EM procedure they adopt in [2]

[4] (2022) Sparse Probabilistic Circuits via Pruning and Growing

<https://arxiv.org/pdf/2211.12551.pdf>

[5] (2020) Random Sum-Product Networks: A Simple and Effective Approach to Probabilistic Deep Learning

<http://proceedings.mlr.press/v115/peharz20a/peharz20a.pdf>

[6] (2020) Einsum Networks: Fast and Scalable Learning of Tractable Probabilistic Circuits

<https://arxiv.org/abs/2004.06231>

<https://github.com/cambridge-mlg/EinsumNetworks>

[7] (2022) Scaling Up Probabilistic Circuits by Latent Variable Distillation

<https://arxiv.org/abs/2210.04398>

[8] (2023) Understanding the Distillation Process from Deep Generative Models to Tractable Probabilistic Circuits

<https://arxiv.org/abs/2302.08086>

Papers highlighted in bold are good starting points.